

Deep Q-Learning With Recurrent Neural Networks

Clare Chen, Vincent Ying, Dillon Laird

Stanford, 2016

Abstract

Deep reinforcement learning models have proven to be successful at learning control policies image inputs. However, they have struggled with learning policies that require longer term information. Recurrent neural network architectures have been used in tasks dealing with longer term dependencies between data points. We investigate these architectures to overcome the difficulties arising from learning policies with long term dependencies.

Introduction

- Recent advances in reinforcement learning have led to human-level or greater performance on a wide variety of games (e.g. Atari 2600 Games).
- Deep Q-networks are limited in that they learn a mapping from a single previous state which consist of a small number of game screens.
- We explore the concepts of deep recurrent Q-networks (DRQN), and also a combination of recurrent neural network (RNN) and deep Q-network (DQN).
- In addition to standard RNN architectures, we also examine augmenting the RNN architecture with an attention mechanism.



Q*bert, Seaquest, and Ms. Pac-Man

Deep Recurrent Q-Learning

- The architecture of DRQN augments DQN's fully connected layer with a LSTM.
- We accomplish this by looking at the last L states:

$$\{s_{t-(L-1)}, \dots, s_t\}$$

- We feed these into a convolution neural network (CNN) to get intermediate outputs and then send those through the RNN:

$$\begin{aligned} \text{CNN}(s_{t-i}) &= x_{t-i} \\ \text{RNN}(x_{t-i}, h_{t-i-1}) &= h_{t-i} \end{aligned}$$

- The final output is used to predict the Q value.

Attention Deep Recurrent Q-Learning

- In **linear attention**, we take the last L hidden states and compute the inner product with learned vector v_a and take a softmax:


$$\{v_a^T h_{t-(L-1)}, \dots, v_a^T h_t\}, a_{t-i} = \text{softmax}(v_a^T h_{t-i})$$

- Then, we take a weighted sum of hidden states to get a context vector c_t that is used to calculate Q :

$$c_t = \sum_{i=0}^{L-1} a_{t-i} h_{t-i}$$

- In **global attention**, we use the current hidden state, h_t , to calculate the attention scores: $\{h_{t-(L-1)}^T h_t, \dots, h_{t-1}^T h_t\}$, $a_{t-i} = \text{softmax}(h_{t-i}^T h_t)$
- We then compute a context vector similar to above and pass it through a final layer before it is used to compute Q :

$$c_t = \sum_{i=1}^{L-1} a_{t-i} h_{t-i}, \tilde{h} = \tanh(W_a[h_t; c_t] + b_a)$$



Architecture of the Attention DRQN

Experiments

- Basic DQN
- Basic RDQN
- RDQN with linear attention.
- RDQN with global attention.

Box

Text or image?

Results

Q^*bert Scores

Algorithms Scores

Random	157
Sarsa	614
Contingency	960
DQN	1952
Human	18900
HNeat Best	1800
HNeat Pixel	1325
DQN Best	4500

Discussion

- Conclusions
- Future work